

# Robust Multiagent Combinatorial Path Finding

## Technical Appendix

### Proofs

In this section, we establish the theoretical guarantees of our approach. We begin by proving the correctness and completeness of the K-Best-Sequencing procedure used in *RCbssEff*, which consists of three steps: (1) reducing the problem of finding a goal sequence allocation  $\gamma$  (an mTSP instance) to an equivalent single-agent E-GTSP instance, (2) enumerating the  $K$  lowest-cost tours using the K-Best-EGTSP algorithm (Alg. 2), and (3) decomposing the  $K^{th}$  tour back into a multi-agent goal sequence allocation  $\gamma$ . Building on these results, we then prove that *RCbssEff*, based on the Robust CBSS framework (Alg. 1), is both complete and optimal.

### Correctness of Step 1: Reducing the Problem of Finding $\gamma$ (mTSP) to an Equivalent Single-Agent E-GTSP Instance

In Step 1, we reduce the problem of finding a goal sequence allocation  $\gamma$  in  $G$  (an mTSP instance) to an equivalent single-agent *E-GTSP* instance represented by  $G'$ . We prove that for any goal sequence allocation  $\gamma$ , there exists a corresponding *E-GTSP* tour with the same cost, ensuring that this reduction preserves both feasibility and cost.

**Notation 1.** Let  $Tour(\gamma)$  denote the *E-GTSP* tour in  $G'$  constructed from a goal sequence allocation  $\gamma$  in  $G$ .

**Lemma 1.** For any feasible goal sequence allocation  $\gamma$  defined over  $G$ , there exists a corresponding *E-GTSP* tour  $Tour(\gamma)$  in  $G'$ .

*Proof.* Let  $\gamma = \langle sq(1), \dots, sq(n) \rangle$  be a feasible allocation in  $G$ , where  $sq(i)$  is the ordered sequence of goals assigned to agent  $i$ .

**Cluster construction.** Each goal  $g_j \in G$  is represented in  $G'$  as a cluster  $C_{g_j}$  containing one vertex per feasible approach orientation, ensuring each goal is represented exactly once with a chosen orientation. Each agent  $i$  is represented as a singleton cluster  $C_0^i$  containing its start position  $s_i$ , ensuring all agents are included.

---

Copyright © 2026, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

**Edge construction.** Type-1 edges in  $G'$  connect agent starts to goals and goals to each other, with costs equal to the shortest-path traversal costs in  $G$ , preserving true movement costs. Zero-cost Type-2 edges connect goals to agent-start clusters and agent-start clusters to each other. These edges allow concatenating all agents' paths into a single *E-GTSP* tour and prevent enforcing a fixed destination goal.

**Tour construction.** We form  $Tour(\gamma)$  in  $G'$  by, for each agent  $i$ , starting at its singleton cluster  $C_0^i$  and visiting the orientation-specific vertex of its assigned goals in the order they appear in  $sq(i)$ , selecting exactly one vertex per cluster. After completing an agent's sequence, a Type-2 edge transitions to the next singleton cluster  $C_0^k$ , concatenating all sequences into a single *E-GTSP* tour.

Since every goal in  $\gamma$  appears exactly once across all agents and every agent's start position is included, the constructed path visits one vertex per cluster in  $G'$  and is thus a valid *E-GTSP* tour. Moreover, no fixed destination cluster is enforced, so the tour accurately reflects the original problem's lack of destination goal constraints.  $\square$

**Lemma 2.** The cost of the *E-GTSP* tour  $Tour(\gamma)$  in  $G'$  equals the cost of the original allocation  $\gamma$  in  $G$ .

*Proof.* The cost of each Type-1 edge between orientation-specific vertices (including edges from agent starts to goals) is inherited directly from the original cost function in  $G$ , which accounts for movement distances and orientations. Zero-cost Type-2 edges do not contribute to the cost. Thus, the total cost of  $Tour(\gamma)$  equals the total cost of executing the allocation  $\gamma$  in  $G$ .  $\square$

**Theorem 3.** Let  $\gamma^*$  be an optimal goal sequence allocation in  $G$ . Then  $Tour(\gamma^*)$  is an optimal tour in  $G'$  for the corresponding *E-GTSP* instance.

*Proof.* By Lemma 1, Step 1 constructs a valid *E-GTSP* tour in  $G'$  for any feasible allocation in  $G$ . By Lemma 2, this construction preserves the cost of the allocation. Therefore, the optimal allocation  $\gamma^*$  corresponds to a tour  $Tour(\gamma^*)$  with the same minimal cost in the *E-GTSP* formulation.  $\square$

## Completeness of Step 2: Enumerating the $K$ Lowest-Cost E-GTSP Tours (Alg. 2)

In Step 2, we enumerate the  $K$  lowest-cost E-GTSP tours in non-decreasing order of cost. We prove that the K-Best-EGTSP algorithm is complete, i.e., it returns all  $K$  lowest-cost tours whenever they exist.

Let  $\mathcal{S}$  denote the set of all feasible tours for a given  $E$ -GTSP instance. A set of tours  $\mathcal{S}_K \subseteq \mathcal{S}$  is a  $K$ -lowest-cost set if (1)  $|\mathcal{S}_K| = K$ , and (2) every tour  $S' \in \mathcal{S} \setminus \mathcal{S}_K$  has cost at least as large as the most expensive tour in  $\mathcal{S}_K$ .

**Theorem 4.** *Given a complete and optimal E-GTSP solver, K-best-EGTSP is guaranteed to return a K-lowest-cost set of tours for any  $K > 0$ .*

*Proof.* Intuitively,  $K$ -best-EGTSP enumerates tours by systematically branching on edges of previously generated tours. Because the underlying  $E$ -GTSP solver is complete and optimal, each new branch eventually explores all feasible tours in non-decreasing order of cost.

Assume, for contradiction, that  $K$ -best-EGTSP fails to return a valid  $K$ -lowest-cost set. Then there exists a returned tour  $s'_j$  such that a cheaper tour  $s_j$  is missing, i.e.,  $\text{cost}(s_j) < \text{cost}(s'_j)$ .

**Case 1:** The first call to  $r\text{EGTSP}$  is made with  $I_c^0 = O_c^0 = \emptyset$ , i.e., directly invoking the base  $E$ -GTSP solver. By assumption, the solver returns the optimal tour. If a cheaper tour  $s_j$  existed, it would have been selected over  $s'_j$ , a contradiction.

**Case 2:** Suppose  $s'_j$  is the  $K^{th}$  selected tour. If  $s_j$  was already in the priority queue when  $s'_j$  was chosen, the best-first ordering would require expanding  $s_j$  first, contradicting the selection of  $s'_j$ . If  $s_j$  was not yet generated when  $s'_j$  was selected, then it must eventually be produced in a later iteration, as  $K$ -best-EGTSP systematically explores all feasible tours by branching on previously generated ones with added inclusion/exclusion constraints. Crucially,  $r\text{EGTSP}$  solves these constrained subproblems by temporarily modifying edge costs: edges required by inclusion constraints receive a large negative offset to enforce their selection, while all other edges either retain their original costs or are penalized. After a tour is found, its cost is re-evaluated using the original edge costs, ensuring that no constrained subproblem can yield a tour with a strictly lower true cost than what could have been found in the unconstrained problem. Therefore, if  $s_j$  were truly cheaper than  $s'_j$ , it would necessarily have been discovered earlier during the enumeration, which contradicts the assumption.

Thus, no such  $s_j$  can exist, proving the algorithm's completeness.  $\square$

## Correctness of Step 3: Decomposing the E-GTSP Tour into a Goal Sequence Allocation

$$\gamma$$

In Step 3, we decompose an E-GTSP tour in  $G'$  into a goal sequence allocation  $\gamma$  in  $G$ . We prove that this decomposition preserves feasibility and cost, yielding an equivalent optimal allocation.

**Notation 2.** Let  $\text{Alloc}(\text{tour})$  denote the goal sequence allocation in  $G$  obtained by decomposing a given  $E$ -GTSP tour in  $G'$ .

**Lemma 5.** *For any  $E$ -GTSP tour in  $G'$ , there exists a corresponding goal sequence allocation  $\gamma$  in  $G$ .*

*Proof.* Let  $\text{tour}$  be an  $E$ -GTSP tour in  $G'$ , represented as an ordered sequence of vertices  $\{s_0, v_{g_{j,o}}, \dots, s_i, \dots, s_n, \dots\}$ , where each  $s_i$  is the initial configuration of agent  $i$ , and each  $v_{g_{j,o}}$  is an orientation-specific copy of goal  $j$  approached from orientation  $o$ . Since each initial agent vertex forms a singleton cluster in  $G'$ , the tour must include exactly  $n$  such agent vertices. To construct  $\text{Alloc}(\text{tour})$ , we partition the tour into  $n$  disjoint segments by cutting at each agent vertex. For each agent  $i$ , let  $sq(i)$  denote the sequence of goal vertices appearing between  $s_i$  and the next agent vertex. Each  $sq(i)$  is assigned to agent  $i$ , preserving the visitation order. Finally, we map each orientation-specific goal vertex  $v_{g_{j,o}} \in sq(i)$  in  $G'$  to its original goal  $v_{g_j}$  in  $G$ , discarding orientation information. This produces a valid goal sequence allocation  $\text{Alloc}(\text{tour})$  in  $G$ .  $\square$

**Lemma 6.** *The cost of the goal sequence allocation in  $G$  is equal to the cost of the original  $E$ -GTSP tour in  $G'$ .*

*Proof.* As described in Lemma 5, the allocation is constructed by partitioning the  $E$ -GTSP tour in  $G'$  into agent-specific segments. Segmentation occurs at Type-2 edges, which are auxiliary and have zero cost. Removing these edges does not affect the total cost. Additionally, mapping orientation-specific goal vertices  $v_{g_{j,o}}$  in  $G'$  to their original vertices  $v_{g_j}$  in  $G$  does not alter the cost: the edge weights between goals in the allocation are inherited from the  $E$ -GTSP traversal, which already accounts for orientation-dependent movement. Thus, the total cost remains unchanged when mapping from  $G'$  back to  $G$ .  $\square$

**Theorem 7.** *Let  $\text{tour}$  be an optimal tour in  $G'$  for the  $E$ -GTSP. Then  $\text{Alloc}(\text{tour})$  is an optimal goal sequence allocation in  $G$ .*

*Proof.* By Lemma 5, Step 3 produces a valid allocation in  $G$  from any  $E$ -GTSP tour in  $G'$ . By Lemma 6, this decomposition preserves the total cost. Thus, if  $\text{tour}$  is optimal for the  $E$ -GTSP in  $G'$ , then  $\text{Alloc}(\text{tour})$  is also optimal for the corresponding goal sequence allocation problem in  $G$ .  $\square$

## Completeness of $RCbssEff$ (Based on the Robust CBSS Framework in Alg. 1)

In this section, we prove that  $RCbssEff$  is *complete*: if a feasible robust solution exists, the algorithm is guaranteed to return one.

**Theorem 8.** *Given the correctness and completeness of Steps 1–3 (which ensure the correctness and completeness of the K-Best-Sequencing procedure),  $RCbssEff$  is solution complete.*

*Proof.* Assume, for contradiction, that  $RCbssEff$  fails to return a solution for some solvable instance. Let  $\pi$  be a feasible robust solution, and let  $\gamma(\pi)$  denote its corresponding goal sequence allocation.

$RCbssEff$  maintains a set  $\mathcal{T}$  of Constraint Trees (CTs), each representing a unique goal sequence allocation, and systematically explores them using CBS to find feasible solutions.

**Case 1:**  $RCbssEff$  never creates a CT tree for  $\gamma(\pi)$ . However,  $RCbssEff$  enumerates all goal sequence allocations using the K-Best-Sequencing method, which invokes *K-best-EGTSP*. Since *K-best-EGTSP* is complete,  $\gamma(\pi)$  must eventually be generated, ensuring that a CT tree for it is created — a contradiction.

**Case 2:** A CT tree  $T$  for  $\gamma(\pi)$  exists, but  $RCbssEff$  never finds  $\pi$  within it. Since  $\pi$  is a feasible robust solution consistent with  $\gamma(\pi)$ , it must correspond to some CT node in  $T$ .  $RCbssEff$  systematically expands all nodes in each CT until a robust solution is found or the tree is fully explored. Thus,  $\pi$  must eventually be reached and verified, contradicting the assumption.

Therefore, no feasible robust solution can be omitted, proving that  $RCbssEff$  is complete.  $\square$

### Optimality of $RCbssEff$ (Based on the Robust CBSS Framework in Alg. 1)

Next, we prove that the solution returned by  $RCbssEff$  is *optimal* with respect to the defined cost function.

**Theorem 9.** *Given the correctness and completeness of Steps 1–3 (which ensure the correctness and completeness of the K-Best-Sequencing procedure), the solution returned by  $RCbssEff$  is optimal.*

*Proof.* Let  $Z$  denote the CT node corresponding to the solution  $\pi$  returned by  $RCbssEff$ . By construction,  $\pi$  consists of tours where each agent visits all its assigned goals. As  $Z$  was expanded, every other CT node  $Z' \in \text{OPEN}$  satisfied  $\text{cost}(Z') \geq \text{cost}(Z)$ .

Assume, for contradiction, that there exists another solution  $\pi'$  with  $\text{cost}(\pi') < \text{cost}(\pi)$ . Let  $\gamma(\pi')$  denote the goal sequence allocation of  $\pi'$ .

**Case 1:** A CT tree for  $\gamma(\pi')$  already exists. Then there must be an unexpanded node  $Z_T$  in this tree with  $\text{cost}(Z_T) \leq \text{cost}(\pi')$ . But since  $RCbssEff$  selected  $Z$  for expansion, we must have  $\text{cost}(Z) \leq \text{cost}(Z_T)$ , which contradicts  $\text{cost}(\pi') < \text{cost}(\pi)$ .

**Case 2:** No CT tree for  $\gamma(\pi')$  has been generated. Then  $\text{cost}(Z) \leq \text{cost}(\gamma(\pi'))$ , where  $\text{cost}(\gamma(\pi'))$  is the minimal cost of its allocation. Since  $\text{cost}(\gamma(\pi')) \leq \text{cost}(\pi')$ , it follows that  $\text{cost}(\pi) \leq \text{cost}(\pi')$ , again contradicting the assumption.

Thus, no cheaper solution exists, and  $RCbssEff$  returns an optimal solution.  $\square$