

Robust Multiagent Combinatorial Path Finding

Technical Appendix

Proofs

In this section, we establish the theoretical guarantees of our approach. We begin by proving the correctness and completeness of the K-Best-Sequencing procedure used in *RCbssEff*, which consists of three steps: (1) reducing the problem of finding a goal sequence allocation γ (an mTSP instance) to an equivalent single-agent E-GTSP instance, (2) enumerating the K lowest-cost tours using the K-Best-EGTSP algorithm (Alg. 2), and (3) decomposing the K^{th} tour back into a multi-agent goal sequence allocation γ . Building on these results, we then prove that *RCbssEff*, based on the Robust CBSS framework (Alg. 1), is both complete and optimal.

Correctness of Step 1: Reducing the Problem of Finding γ (mTSP) to an Equivalent Single-Agent E-GTSP Instance

In Step 1, we reduce the problem of finding a goal sequence allocation γ in G (an mTSP instance) to an equivalent single-agent *E-GTSP* instance represented by G' . We prove that for any goal sequence allocation γ , there exists a corresponding *E-GTSP* tour with the same cost, ensuring that this reduction preserves both feasibility and cost.

Notation 1. Let $Tour(\gamma)$ denote the *E-GTSP* tour in G' constructed from a goal sequence allocation γ in G .

Lemma 1. For any feasible goal sequence allocation γ defined over G , there exists a corresponding *E-GTSP* tour $Tour(\gamma)$ in G' .

Proof. Let $\gamma = \langle sq(1), \dots, sq(n) \rangle$ be a feasible allocation in G , where $sq(i)$ is the ordered sequence of goals assigned to agent i .

Cluster construction. Each goal vertex $v_{g_j} \in V_g$ is represented in G' as a cluster C_{g_j} containing one vertex per feasible approach orientation, ensuring each goal is represented exactly once with a chosen orientation. Each agent i is represented as a singleton cluster C_0^i containing its start position s_i , ensuring all agents are included.

Copyright © 2026, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

Edge construction. Type-1 edges in G' connect agent starts to goals and goals to each other, with costs equal to the shortest-path traversal costs in G , preserving true movement costs. Zero-cost Type-2 edges connect goals to agent-start clusters and agent-start clusters to each other. These edges allow concatenating all agents' paths into a single *E-GTSP* tour and prevent enforcing a fixed destination goal.

Tour construction. We form $Tour(\gamma)$ in G' by, for each agent i , starting at its singleton cluster C_0^i and visiting the orientation-specific vertex of its assigned goals in the order they appear in $sq(i)$, selecting exactly one vertex per cluster. After completing an agent's sequence, a Type-2 edge transitions to the next singleton cluster C_0^k , concatenating all sequences into a single *E-GTSP* tour.

Since every goal in γ appears exactly once across all agents and every agent's start position is included, the constructed path visits one vertex per cluster in G' and is thus a valid *E-GTSP* tour. Moreover, no fixed destination cluster is enforced, so the tour accurately reflects the original problem's lack of destination goal constraints. \square

Lemma 2. The cost of the *E-GTSP* tour $Tour(\gamma)$ in G' equals the cost of the original allocation γ in G .

Proof. The cost of each Type-1 edge between orientation-specific vertices (including edges from agent starts to goals) is inherited directly from the original cost function in G , which accounts for movement distances and orientations. Zero-cost Type-2 edges do not contribute to the cost. Thus, the total cost of $Tour(\gamma)$ equals the total cost of executing the allocation γ in G . \square

Theorem 3. Let γ^* be an optimal goal sequence allocation in G . Then $Tour(\gamma^*)$ is an optimal tour in G' for the corresponding *E-GTSP* instance.

Proof. By Lemma 1, Step 1 constructs a valid *E-GTSP* tour in G' for any feasible allocation in G . By Lemma 2, this construction preserves the cost of the allocation. Therefore, the optimal allocation γ^* corresponds to a tour $Tour(\gamma^*)$ with the same minimal cost in the *E-GTSP* formulation. \square

Completeness of Step 2: Enumerating the K Lowest-Cost E-GTSP Tours (Alg. 2)

In Step 2, we enumerate the K lowest-cost E-GTSP tours in non-decreasing order of cost. We prove that the K-Best-EGTSP algorithm is complete, i.e., it returns all K lowest-cost tours whenever they exist.

Let \mathcal{S} denote the set of all feasible tours for a given E -GTSP instance. A set of tours $\mathcal{S}_K \subseteq \mathcal{S}$ is a K -lowest-cost set if (1) $|\mathcal{S}_K| = K$, and (2) every tour $S' \in \mathcal{S} \setminus \mathcal{S}_K$ has cost at least as large as the most expensive tour in \mathcal{S}_K .

Theorem 4. *Given a complete and optimal E-GTSP solver, K-best-EGTSP is guaranteed to return a K-lowest-cost set of tours for any $K > 0$.*

Proof. Intuitively, K -best-EGTSP enumerates tours by systematically branching on edges of previously generated tours. Because the underlying E -GTSP solver is complete and optimal, each new branch eventually explores all feasible tours in non-decreasing order of cost.

Assume, for contradiction, that K -best-EGTSP fails to return a valid K -lowest-cost set. Then there exists a returned tour s'_j such that a cheaper tour s_j is missing, i.e., $\text{cost}(s_j) < \text{cost}(s'_j)$.

Case 1: The first call to $r\text{EGTSP}$ is made with $I_c^0 = O_c^0 = \emptyset$, i.e., directly invoking the base E -GTSP solver. By assumption, the solver returns the optimal tour. If a cheaper tour s_j existed, it would have been selected over s'_j , a contradiction.

Case 2: Suppose s'_j is the K^{th} selected tour. If s_j was already in the priority queue when s'_j was chosen, the best-first ordering would require expanding s_j first, contradicting the selection of s'_j . If s_j was not yet generated when s'_j was selected, then it must eventually be produced in a later iteration, as K -best-EGTSP systematically explores all feasible tours by branching on previously generated ones with added inclusion/exclusion constraints. Crucially, $r\text{EGTSP}$ solves these constrained subproblems by temporarily modifying edge costs: edges required by inclusion constraints receive a large negative offset to enforce their selection, while all other edges either retain their original costs or are penalized. After a tour is found, its cost is re-evaluated using the original edge costs, ensuring that no constrained subproblem can yield a tour with a strictly lower true cost than what could have been found in the unconstrained problem. Therefore, if s_j were truly cheaper than s'_j , it would necessarily have been discovered earlier during the enumeration, which contradicts the assumption.

Thus, no such s_j can exist, proving the algorithm's completeness. \square

Correctness of Step 3: Decomposing the E-GTSP Tour into a Goal Sequence Allocation

$$\gamma$$

In Step 3, we decompose an E-GTSP tour in G' into a goal sequence allocation γ in G . We prove that this decomposition preserves feasibility and cost, yielding an equivalent optimal allocation.

Notation 2. Let $\text{Alloc}(\text{tour})$ denote the goal sequence allocation in G obtained by decomposing a given E -GTSP tour in G' .

Lemma 5. *For any E -GTSP tour in G' , there exists a corresponding goal sequence allocation γ in G .*

Proof. Let tour be an E -GTSP tour in G' , represented as an ordered sequence of vertices $\{s_0, v_{g_{j,o}}, \dots, s_i, \dots, s_n, \dots\}$, where each s_i is the initial configuration of agent i , and each $v_{g_{j,o}}$ is an orientation-specific copy of goal j approached from orientation o . Since each initial agent vertex forms a singleton cluster in G' , the tour must include exactly n such agent vertices. To construct $\text{Alloc}(\text{tour})$, we partition the tour into n disjoint segments by cutting at each agent vertex. For each agent i , let $sq(i)$ denote the sequence of goal vertices appearing between s_i and the next agent vertex. Each $sq(i)$ is assigned to agent i , preserving the visitation order. Finally, we map each orientation-specific goal vertex $v_{g_{j,o}} \in sq(i)$ in G' to its original goal v_{g_j} in G , discarding orientation information. This produces a valid goal sequence allocation $\text{Alloc}(\text{tour})$ in G . \square

Lemma 6. *The cost of the goal sequence allocation in G is equal to the cost of the original E -GTSP tour in G' .*

Proof. As described in Lemma 5, the allocation is constructed by partitioning the E -GTSP tour in G' into agent-specific segments. Segmentation occurs at Type-2 edges, which are auxiliary and have zero cost. Removing these edges does not affect the total cost. Additionally, mapping orientation-specific goal vertices $v_{g_{j,o}}$ in G' to their original vertices v_{g_j} in G does not alter the cost: the edge weights between goals in the allocation are inherited from the E -GTSP traversal, which already accounts for orientation-dependent movement. Thus, the total cost remains unchanged when mapping from G' back to G . \square

Theorem 7. *Let tour be an optimal tour in G' for the E -GTSP. Then $\text{Alloc}(\text{tour})$ is an optimal goal sequence allocation in G .*

Proof. By Lemma 5, Step 3 produces a valid allocation in G from any E -GTSP tour in G' . By Lemma 6, this decomposition preserves the total cost. Thus, if tour is optimal for the E -GTSP in G' , then $\text{Alloc}(\text{tour})$ is also optimal for the corresponding goal sequence allocation problem in G . \square

Completeness of $RCbssEff$ (Based on the Robust CBSS Framework in Alg. 1)

In this section, we prove that $RCbssEff$ is *complete*: if a feasible robust solution exists, the algorithm is guaranteed to return one.

Theorem 8. *Given the correctness and completeness of Steps 1–3 (which ensure the correctness and completeness of the K-Best-Sequencing procedure), $RCbssEff$ is solution complete.*

Proof. Assume, for contradiction, that $RCbssEff$ fails to return a solution for some solvable instance. Let π be a feasible robust solution, and let $\gamma(\pi)$ denote its corresponding goal sequence allocation.

$RCbssEff$ maintains a set \mathcal{T} of Constraint Trees (CTs), each representing a unique goal sequence allocation, and systematically explores them using CBS to find feasible solutions.

Case 1: $RCbssEff$ never creates a CT tree for $\gamma(\pi)$. However, $RCbssEff$ enumerates all goal sequence allocations using the K-Best-Sequencing method, which invokes *K-best-EGTSP*. Since *K-best-EGTSP* is complete, $\gamma(\pi)$ must eventually be generated, ensuring that a CT tree for it is created — a contradiction.

Case 2: A CT tree T for $\gamma(\pi)$ exists, but $RCbssEff$ never finds π within it. Since π is a feasible robust solution consistent with $\gamma(\pi)$, it must correspond to some CT node in T . $RCbssEff$ systematically expands all nodes in each CT until a robust solution is found or the tree is fully explored. Thus, π must eventually be reached and verified, contradicting the assumption.

Therefore, no feasible robust solution can be omitted, proving that $RCbssEff$ is complete. \square

Optimality of $RCbssEff$ (Based on the Robust CBSS Framework in Alg. 1)

Next, we prove that the solution returned by $RCbssEff$ is *optimal* with respect to the defined cost function.

Theorem 9. *Given the correctness and completeness of Steps 1–3 (which ensure the correctness and completeness of the K-Best-Sequencing procedure), the solution returned by $RCbssEff$ is optimal.*

Proof. Let Z denote the CT node corresponding to the solution π returned by $RCbssEff$. By construction, π consists of tours where each agent visits all its assigned goals. As Z was expanded, every other CT node $Z' \in \text{OPEN}$ satisfied $\text{cost}(Z') \geq \text{cost}(Z)$.

Assume, for contradiction, that there exists another solution π' with $\text{cost}(\pi') < \text{cost}(\pi)$. Let $\gamma(\pi')$ denote the goal sequence allocation of π' .

Case 1: A CT tree for $\gamma(\pi')$ already exists. Then there must be an unexpanded node Z_T in this tree with $\text{cost}(Z_T) \leq \text{cost}(\pi')$. But since $RCbssEff$ selected Z for expansion, we must have $\text{cost}(Z) \leq \text{cost}(Z_T)$, which contradicts $\text{cost}(\pi') < \text{cost}(\pi)$.

Case 2: No CT tree for $\gamma(\pi')$ has been generated. Then $\text{cost}(Z) \leq \text{cost}(\gamma(\pi'))$, where $\text{cost}(\gamma(\pi'))$ is the minimal cost of its allocation. Since $\text{cost}(\gamma(\pi')) \leq \text{cost}(\pi')$, it follows that $\text{cost}(\pi) \leq \text{cost}(\pi')$, again contradicting the assumption.

Thus, no cheaper solution exists, and $RCbssEff$ returns an optimal solution. \square